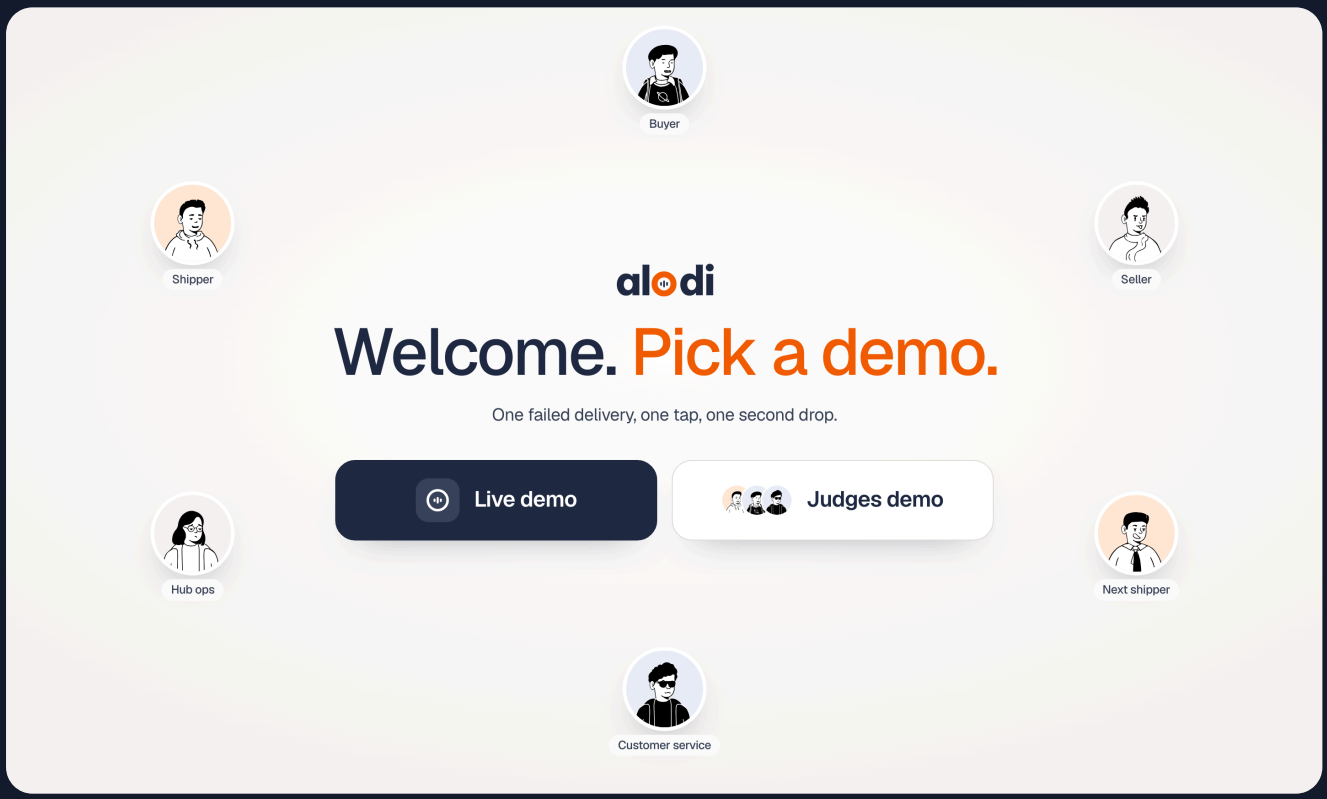
alodi

Make the second drop land.

Alodi is the AI agent that rescues failed deliveries for Vietnamese e-commerce. The shipper taps once. Alodi reaches the buyer by chat and voice, and the next attempt arrives knowing when, where and who.



TEAM	Develarper
EVENT	Sea × OpenAI Regional Codex Hackathon, Ho Chi Minh City, 31 October 2026
BUILD DIRECTION	Autonomous & Adaptive AI · Deep Domain AI

1 The problem

Delivery scheduling happens outside the system, and the failure context is lost with it.

Shopee does not let buyers select a delivery window, so scheduling depends on a phone call between buyer and shipper. When that call fails, the reason the delivery could not be completed is reduced to a single status code. The next attempt runs under the same conditions, and the return cost falls on the seller.

₫40,000	8–12%	0	≈ 5
return fee per failed order, paid by the seller	of orders refused at the door, nationwide	delivery windows a Shopee buyer can choose	full-time callers an area needs to phone 320 buyers twice a day

Survey of buyers and shippers

Placeholder values until the interview export lands

86%	66%	27%	83%
of buyers want to choose a delivery window	of buyers called back and couldn't reach the shipper	of shippers' calls to buyers go unanswered	of shippers would use an Alodi-like system

100 buyers · 12 shippers from 6 carriers

₫473M lost a month in one area of 20 shippers and 4,000 orders a day



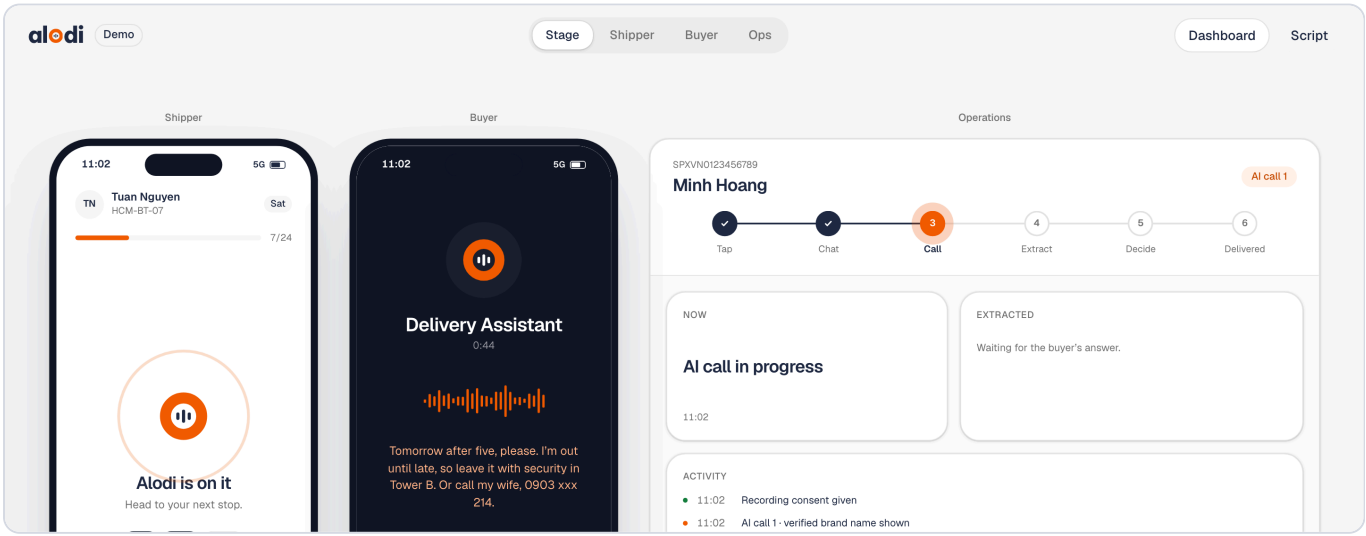
● Return fees ₫288M ● Redelivery pay ₫101M ● CS call-backs ₫45M ● Door time ₫39M

Source: Shopee Return & Refund Policy §7.1; GHN; VECOM via Nguoi Lao Dong, 19 Jul 2025; Shopee Help Center 79559; team cost model (modelled).

2 The solution

One tap in. One context card out.

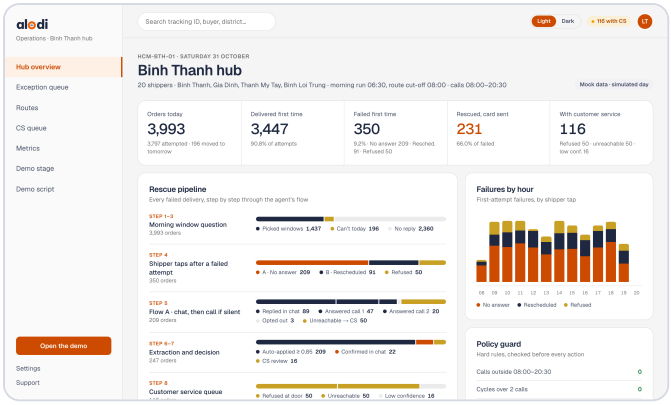
01	02	03	04	05
Tap	**Reach**	**Understand**	**Decide**	**Redeliver**
No answer, Rescheduled or Refused	Chat at once; a voice call if silent, 08:00–20:30 only	The answer becomes when, where and who	Clear answers apply; the rest goes to a person	The next attempt carries the card



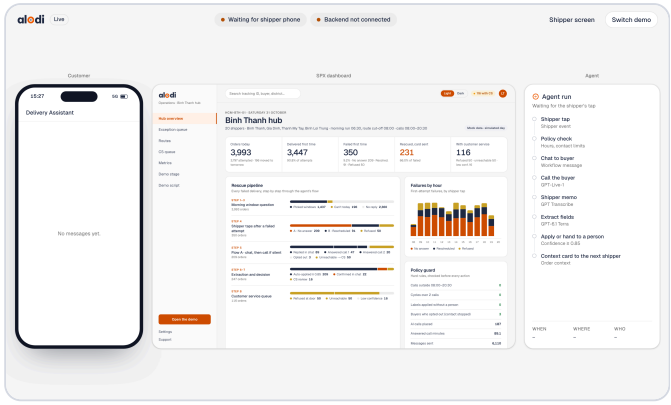
The rescue from the judges' demo: shipper's app, buyer's call and operations, in sync.

3 The product

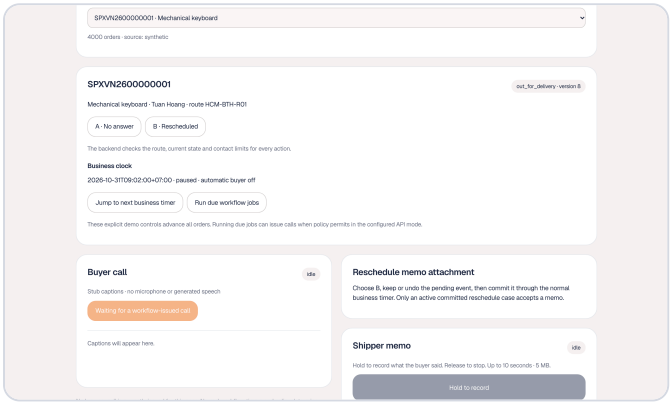
Three working surfaces on one rescue workflow.



Hub dashboard: pipeline by business step, policy guard, exception queue, routes.



Live stage: the shipper uses a real phone; customer, dashboard and agent run sync here.



Backend workspace: real orders, business timers, calls issued by policy, and the order's evidence.

Hard rules, enforced in code

- Calls only 08:00–20:30
- At most two calls a cycle
- Applies alone only at confidence ≥ 0.85
- "Stop" ends all contact
- Shippers can undo a tap for 2 minutes
- A person approves every label and return

COMPONENT	WHAT RUNS TODAY
Workflow backend	FastAPI and SQLite with migrations: morning availability, shipper outcomes, rescue policy, customer-service decisions and a durable scheduler that survives restarts, on 4,000 synthetic orders. **193 automated tests pass**, with AI providers simulated, and CI checks every change.
End-to-end path	On the running backend: a No answer tap opens a rescue case, the 60-minute timer fires, and the policy-checked call is issued.
Voice and memo	GPT-Live-1 over WebRTC, GPT Transcribe and GPT-6.1 Terra extraction, with consent gates, a 180-second call cap and memo limits. Stub and live modes.
Interfaces	Landing, a judges' demo with a scripted scenario, an SPX-style dashboard on mock operational data, and the live stage with its shipper phone screen, whose backend connection is the next step.

4 What is different

Context, not routing

Route optimisers fix the path. The failure sits in a phone call. Alodi captures what that call loses and gives it to the next attempt.

The tap is the trigger

ClickPost and Shipsy start from a carrier NDR code. Alodi starts from a shipper's one meaningful tap, so each path gets the right response.

A hard line in code

The agent asks and calls. Refusals, labels and returns go to a person. Hours, contact limits and the 0.85 threshold sit in code, not in a prompt.

5 OpenAI at the core

MODEL OR API	JOB IN PRODUCTION	ENDPOINT	STATUS
GPT-Live-1	Natural voice calls that listen while speaking and hand reasoning to the agent	v1/live/sessions	In this build
GPT-6.1 Terra	Orchestrates each case under policy and writes the card with Structured Outputs	v1/responses	In this build
GPT Transcribe	Turns the shipper's 10-second memo into text	v1/audio/transcriptions	In this build
Decisions API on GPT-6 Luna	Triages every message cheaply: which case, which window, how sure	v1/decisions	Production
Agents SDK and Evals	Tools, handoffs and traces; held-out evals gate every prompt or model change	Responses API	Production
gpt-oss-20b	Offline fallback on the area server when the network drops	Local	Production

The OpenAI API costs about ₫7.2M a month for one area in our cost model. Names, phones and house numbers are masked before text reaches a model.

6 Who benefits, and how much

SAVED A MONTH, ONE AREA · PROJECTED

₫187M

Lost today **₫473M**

With Alodi **₫286M**

40% of the loss, after ₫27M running cost, at a best-in-class rescue rate.

FROM ANSWER TO CARD · TARGET

< 10 s

0 s — 10 s

- **Shippers** skip repeat trips
- **Buyers** are asked, not chased
- **Sellers** pay fewer return fees
- **Carriers** scale without callers

Source: Team cost model, numbers PDF; Shipsy rescue range (vendor claim). Not yet measured in a pilot.

7 Build direction

Autonomous & Adaptive AI: an operations agent runs a rescue without a dispatcher, adapts to silence, restarts and unclear answers, and escalates when it should not act alone.

Deep Domain AI: shipper outcomes, route cut-offs, Shopee's 2–3 contacts in 5 days and calling hours shape every decision.

Next: full build Nov 2026 · shadow mode at one area Q1 2027 · real channels Q2 · region H2 · national 2028.

8 How we used Codex

Draft. The team confirms this section against what actually happened, then removes this line.

- **Specify:** PRD and rules into scripts and acceptance checks
- **Build:** pair-programming on workflow, voice and interface code
- **Review:** pull requests checked by lint, build and test CI; a person merges
- **Evaluate:** one change a round, kept only if held-out F1 rises, target ≥ 0.90

Sources

SOURCE	WHAT WE USE	WHERE
Shopee Return & Refund Policy §7.1	Seller pays return shipping on failed orders	help.shopee.vn/portal/4/article/77251
GHN, who pays return shipping	₫40,000 cap per order since 1 Apr 2025	ghn.vn/blogs/tip-ban-hang/tra-hang-shopee-ai-chiu-phi
Nguoi Lao Dong, 19 Jul 2025	VECOM: 8–12% of orders refused nationwide	tuoitre.vn/nld/nhuc-nhoi-nan-bom-hang-19625071822023801.htm
Shopee Help Center 79559	No specific delivery time can be chosen	help.shopee.vn/portal/4/article/79559
Shopee Help Center 79569	2–3 contact attempts, reschedule within 5 days	help.shopee.vn/portal/4/article/79569
GHN, shipper pay per order	₫6,000–8,000 per delivered order	ghn.vn/blogs/tip-ban-hang/mot-don-hang-shipper-shopee-duoc-bao-nhieu-tien
Shipsy, NDR management	Rescue 20–30% industry, 50–65% best-in-class (vendor claim)	shipsy.ai/insights/what-is-ndr-non-delivery-report-management
OpenAI, GPT-Live-1	v1/live/sessions; $0.05 a minute	developers.openai.com/api/docs/models/gpt-live-1
OpenAI, GPT-6.1 Terra	Responses API; $1 / $5 per 1M tokens	developers.openai.com/api/docs/models/gpt-6.1-terra
OpenAI, Decisions API	v1/decisions on gpt-6-luna; $0.10 per 1M input tokens	openai.com/index/decisions-api
Team cost model	One area: every input, formula and result	Numbers PDF, submitted with this write-up
Alodi project report	Full design, data, privacy, competitors and roadmap	Report PDF, submitted with this write-up
Competitor comparison	ClickPost, Shipsy and others, from public product pages	Report PDF, chapter 11
Notionists by Zoish	Illustrated people in the product (CC0)	dicebear.com/styles/notionists

How to read our numbers. Sourced figures come from the public documents above, checked 31 October 2026. Survey values are placeholders until the interview export lands. The monthly loss is modelled from them with stated assumptions. Savings are projected until a pilot measures them, and ten seconds is a target. Dashboard figures in the screenshots are simulated.

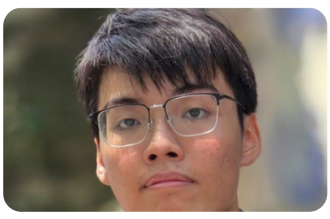
Team Develarper



Nomura Yoshio
Team Lead · Presenter



Long Quan Ton
AI Engineer · Solution Architect



Khanh Tuong Huynh
AI Engineer · Deployment



Thanh Loi Tran
Data Engineer · Backend